

Dimensionality Reduction

Machine Learning II

Master in Business Analytics and Big Data

acastellanos@faculty.ie.edu


Reduce the dimension of the data set → **smaller** data set

...to capture *most of the interesting behavior* of the data


feature selection
or
dimensionality reduction.

Why care about high dimensions?


1-D: 42% of data captured.




2-D: 14% of data captured.



3-D: 7% of data captured.



4-D: 3% of data captured.




Why care about high dimensions?

- **Traditional** statistical techniques apply to problems where the **number of observations (n) is much greater than the number of features (p)**.
- **In recent years, technology changed the game**
 - Example: Understanding people's online shopping patterns considers as *features* all the search terms and interactions with the web by every user. In this case:

$$\begin{aligned}n &\approx 1000 \\ p &\gg n\end{aligned}$$

Principal Component Analysis



Source:
Wikipedia Commons

PCA (2)

- PCA is a technique for **reducing the dimension** of a dataset
- This method seeks to find a **smaller number of representative variables** (**linear combinations of the predictors**, known as principal components PCs), which capture the most possible variance.
- **The first PC captures the most variability of all possible linear combinations.**
 - Then, subsequent PCs are derived such that these linear combinations capture the most remaining variability while also being uncorrelated with all previous PCs.

PCA (3)

- PCA works with the correlation between variables.
- **If the variables are *uncorrelated*, there is no point in computing PCA!**
- PCA is a **technique** (not an algorithm). It is absolutely *exploratory* or *preparatory*.
 - With the appropriate tools it becomes almost ridiculously easy to perform.
 - It's important to **understand the basics** behind to ensure that this method is applied properly.

Advantages & Caveats

- It creates components that are **uncorrelated**.
- PCA **does not consider the response variable** when summarizing variability. It is **unsupervised**.
 - For a supervised equivalent, use “LDA/QDA”.
- Sensitive to **scales** and (**skewed**) distributions
 - PCA focuses on **identifying the data structure** based on measurement scales rather than the important relationships within the data.

Scaling

```
# load libraries
library(caret)
# load the dataset
data(iris)
# summarize data
summary(iris[,1:4])
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100
1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300
Median :5.800	Median :3.000	Median :4.350	Median :1.300
Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199
3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800
Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500

Approach to standardization (**scaling and centering**) using CARET package



```
# calculate the pre-process parameters from the dataset
preprocessParams <- preProcess(iris[,1:4], method=c("center", "scale"))
# transform the dataset using the parameters
transformed <- predict(preprocessParams, iris[,1:4])
# summarize the transformed dataset
summary(transformed)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
Min. :-1.86378	Min. :-2.4258	Min. :-1.5623	Min. :-1.4422
1st Qu.: -0.89767	1st Qu.: -0.5904	1st Qu.: -1.2225	1st Qu.: -1.1799
Median : -0.05233	Median : -0.1315	Median : 0.3354	Median : 0.1321
Mean : 0.00000	Mean : 0.0000	Mean : 0.0000	Mean : 0.0000
3rd Qu.: 0.67225	3rd Qu.: 0.5567	3rd Qu.: 0.7602	3rd Qu.: 0.7880
Max. : 2.48370	Max. : 3.0805	Max. : 1.7799	Max. : 1.7064

Skewness

```
# load libraries
library(mlbench)
library(caret)
# load the dataset
data(PimaIndiansDiabetes)
# summarize pedigree and age
summary(PimaIndiansDiabetes[,7:8])
```

pedigree		age	
Min.	:0.0780	Min.	:21.00
1st Qu.:	:0.2437	1st Qu.:	:24.00
Median	:0.3725	Median	:29.00
Mean	:0.4719	Mean	:33.24
3rd Qu.:	:0.6262	3rd Qu.:	:41.00
Max.	:2.4200	Max.	:81.00



Approach to fixing
skewness, using CARET
package

```
# calculate the pre-process parameters from the dataset
preprocessParams <- preProcess(PimaIndiansDiabetes[,7:8], method=c("BoxCox"))
# transform the dataset using the parameters
transformed <- predict(preprocessParams, PimaIndiansDiabetes[,7:8])
# summarize the transformed dataset (note pedigree and age)
summary(transformed)
```

pedigree		age	
Min.	: -2.5510	Min.	: 0.8772
1st Qu.:	: -1.4116	1st Qu.:	: 0.8815
Median	: -0.9875	Median	: 0.8867
Mean	: -0.9599	Mean	: 0.8874
3rd Qu.:	: -0.4680	3rd Qu.:	: 0.8938
Max.	: 0.8838	Max.	: 0.9019

How?

```
> Result <- prcomp( data )
```

Summary Stats

Measures of
LOCATION

Mean
Median
Percentiles

Measures of
S P R E A D


Variance
Standard Deviation
Interquartile Range

Measures of
ASSOCIATION

Covariance
Correlation

Mean, median and percentiles


- **Mean** is the average
- **Median** is the middle value in the sorted list of values.
- **Percentiles**: 25th, 50th and 75th percentiles respectively
 - *value below which a given percentage of observations in a group of observations fall*




Mean or Median?

- **Median is preferred as, mean is unstable, and whenever there's skewness, median is more stable.**


Right skewed



Symmetric



Left skewed



**Most typical:
prices, salaries, etc.**

YAHOO!
FINANCE

Recent
Quotes you view appear here for
quick access.

Quote Lookup

Finance Home

Search Finance

Search Web

Wed, Dec 31, 2014, 1:35pm EST - US Markets close in 2 hrs and 25 mins

U.S. Median Home Price in November
Increases 15 Percent From a Year Ago,


Variance, Stdev and Interquartile range

- **Variance** (n is sample size)

Always positive, measures *if the points are very close to each other* and to the mean (small), or spread out (high).

$$\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n - 1}$$

- **Standard deviation**: Square root of the variance. It's **on the scale of the original data**. Represents the average distance from the mean.
- The **interquartile range** is the 3rd quartile minus the 1st quartile.



Covariance & Correlation

- **Covariance** between x and y

Measures how two variables change together

Positive: Y increases as X increases

Negative: Y decreases as X increases

WARN: Range is the same as the variables.

$$\frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{n - 1}$$

- **Correlation** (coefficient of) is covariance divided by the product of the two standard deviations

- Ranges between -1 and 1.

Theory

- **Covariance Matrix**

$$\Sigma = \begin{pmatrix} \text{cov}(x, x) & \text{cov}(x, y) & \cdots \\ \text{cov}(y, x) & \text{cov}(y, y) & \\ \vdots & & \ddots \end{pmatrix}$$

- Because of the symmetric nature of correlation (and covariance), **this matrix is equals to its transpose** (reflects elements along its diagonal).

Reminder on Matrix Operations

Inverse Matrix

$$A = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \rightarrow A^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

Transpose Matrix

$$\begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{pmatrix}^T = \begin{pmatrix} a_{11} & a_{21} & a_{31} \\ a_{12} & a_{22} & a_{32} \\ a_{13} & a_{23} & a_{33} \end{pmatrix}$$

Theory (2)

- So,... if we invoke the ***spectral decomposition theorem*** at this point, we find that
 - For any symmetric $N \times N$ matrix A , there exists another orthogonal* matrix U , such that:

$$B = \begin{pmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_N \end{pmatrix} = U^{-1} \cdot A \cdot U$$

- λ_i are the **eigenvalues**
- Columns of U , are the **eigenvectors** of A

*Orthogonal: $U^T = U^{-1}$; $U^T U = U U^T = 1$


What is an eigenvector?

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 1 \\ 3 \end{pmatrix} = \begin{pmatrix} 11 \\ 5 \end{pmatrix}$$


$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$

In the first line,
the resulting vector is not an integer multiple of the original vector,
whereas in the second line, the example is exactly **4** times the vector we began
with.

Why is this?



$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$



$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix}$ Is a **transformation matrix**. If you multiply this matrix on the left of a vector, the answer is another vector that is transformed from it's original position.

Eigenvectors

Eigenvectors (of a given matrix) are vectors that suffer no transformation, just scaling, when multiplied.

$\begin{pmatrix} 3 \\ 2 \end{pmatrix}$ is an **eigenvector** of the *transformation matrix* $\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix}$.

Properties

- 1) Eigenvectors can only be found for **square matrices**. Not every square matrix has eigenvectors. $n \times n$ matrices has (if any) n eigenvectors.
- 2) If vector is **scaled** by some amount before transforming it, the result is a multiple of it, too.

$$2 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 6 \\ 4 \end{pmatrix}$$

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 6 \\ 4 \end{pmatrix} = \begin{pmatrix} 24 \\ 16 \end{pmatrix} = 4 \times \begin{pmatrix} 6 \\ 4 \end{pmatrix}$$


- 3) Eigenvectors are perpendicular (**orthogonal**). That means that we express data in terms of the eigenvectors instead of the x and y axes.
- 4) Due to their scaling property (2), we're interested only in eigenvectors of **length = 1**.

Eigenvalues

Remember the amount by which the original vector was scaled after multiplication by the square matrix?

Remember that amount was always the same?

That is the **eigenvalue**.

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$


No matter what multiple of the eigenvector we took before we multiplied it by the square matrix, we would always get 4 times the scaled vector as our result



Eigenvalues




Why is this important?

- We can change the variables in our observations from a square matrix A into a diagonal matrix, B , which contains the same information.
- The purpose is to **find the directions** in a new coordinate system, **where is more suitable to describe the data** than was in the original coordinate system.

Example



Distribution of the two variables along x and y .



Distribution of the two variables along x' and y' .


Original pictures taken from “Data analysis with open source tools”

Interpretation

- ***Eigenvalues*** point along the directions of greatest variance.
 - The eigenvalue (corresponding to each eigenvector) is a measure of the **width of the distribution** along this direction.
 - Variables measured along the principal directions are **uncorrelated** with each other (*take a look to the correlation matrix*).
- If the data points are distributed as a globular cloud, ***eigenvectors*** will give us the directions of the principal axes of the cloud of data points and the ***eigenvalues*** will give us the length of the cloud along each of these directions

Why is this useful?


- The PCA uses the information contained in the mutual correlations between variables to **identify** those that are **redundant**.
- The **irrelevant** variables are those corresponding to *small eigenvalues*



Example

- Let's produce some sample data with a strong correlation.
- 2 variables (dimensions), X and Y .
- Y is simply, $X \pm$ noise between 0 and 2.

```
x <- runif(100, 0, 16)
for(i in seq(along=x)) {
  y[i] = x[i] + (randSign()*runif(1, 0, 2))
}
data <- data.frame(x,y)
```



PCA on sample data

```
> pca <- prcomp(data)
> print(pca)
Standard deviations:
[1] 6.9073683 0.7656859

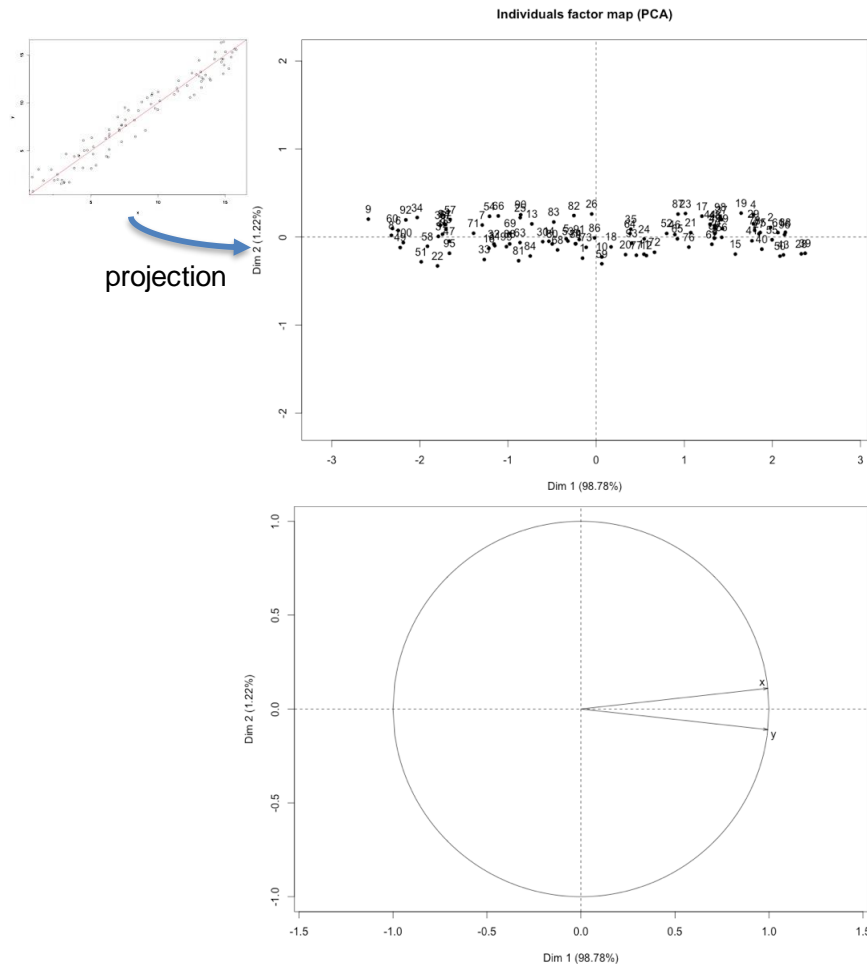
Rotation:
      PC1      PC2
x 0.6959660 -0.7180748
y 0.7180748  0.6959660

> summary(pca)
Importance of components:
               PC1      PC2
Standard deviation   6.9074 0.76569
Proportion of Variance 0.9879 0.01214
Cumulative Proportion 0.9879 1.00000

> pca$sd^2
[1] 47.7117365  0.5862748
```

- **PC1 stdev is much greater than PC2**, which means that **PC1 explains most of the variance**.
- In the summary, you can see how **PC1 dimension explains 98,79%** of the variance.
- **Rotation is the matrix of eigenvectors**. PC1 dimension presents almost the same value for X and Y , which means that they're correlated.
- And finally, just for curiosity, what you can check is the eigenvalues themselves, where λ_1 is a lot bigger than λ_2 , explaining its importance.

And PCA, graphically



The **projection** of the variables onto the 2 principal components result in the expected stretched ellipsoid.

- X and Y , are parallel to the first PC \rightarrow **they are important variables.**
- Both vectors are almost identical (parallel) \rightarrow **they're strongly correlated.**
- Most important take away is that X and Y are redundant.

A more complex example

Wine quality dataset: <http://archive.ics.uci.edu/ml/datasets/Wine+Quality>

- 1 - fixed acidity
- 2 - volatile acidity
- 3 - citric acid
- 4 - residual sugar
- 5 - chlorides
- 6 - free sulfur dioxide
- 7 - total sulfur dioxide
- 8 - density
- 9 - pH
- 10 - sulphates
- 11 - alcohol
- 12 - quality (score between 0 and 10)

```
wine <- read.csv( "winequality-white.csv",  
sep=';', header=TRUE )  
pc <- prcomp( wine )  
plot( pc )
```

Same example, explained in: <http://blog.haunschmid.name/dimensionality-reduction-1-understanding-pca-and-ica-using-r/>

Exercise

Do the values in dataset range within the same scales?

What variables are more representative?

How many variables do we need to explain, at least, 75% of the variance?

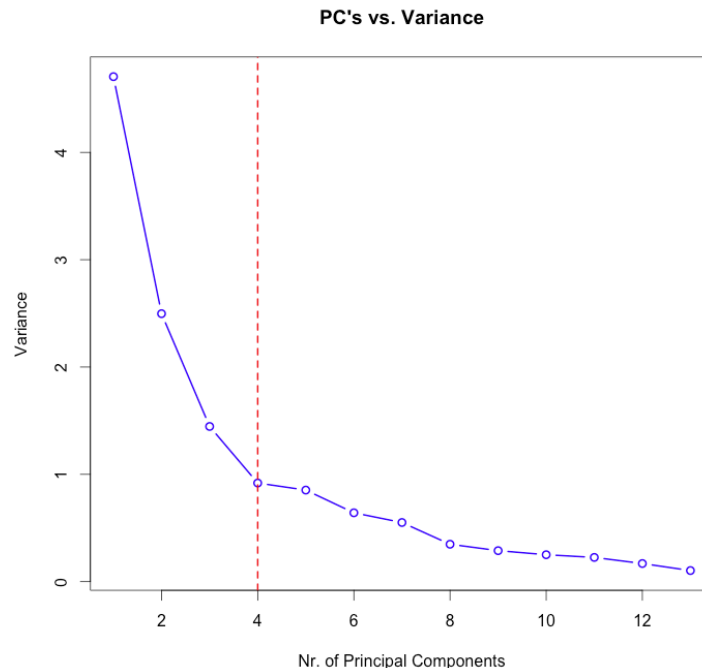
Ranges

```
> head(wine)
  V1    V2    V3    V4    V5    V6    V7    V8    V9    V10    V11    V12    V13    V14
1  1 14.23  1.71  2.43 15.6 127  2.80  3.06  0.28  2.29  5.64  1.04  3.92 1065
2  1 13.20  1.78  2.14 11.2 100  2.65  2.76  0.26  1.28  4.38  1.05  3.40 1050
3  1 13.16  2.36  2.67 18.6 101  2.80  3.24  0.30  2.81  5.68  1.03  3.17 1185
4  1 14.37  1.95  2.50 16.8 113  3.85  3.49  0.24  2.18  7.80  0.86  3.45 1480
5  1 13.24  2.59  2.87 21.0 118  2.80  2.69  0.39  1.82  4.32  1.04  2.93  735
6  1 14.20  1.76  2.45 15.2 112  3.27  3.39  0.34  1.97  6.75  1.05  2.85 1450

> std_wine <- as.data.frame(scale(wine[2:14]))
  V2    V3    V4    V5    V6    V7    V8    V9    V10    V11    V12    V13    V14
1 1.51 -0.56  0.23 -1.17  1.91  0.81  1.03 -0.66  1.22  0.25  0.36  1.84  1.01
2 0.25 -0.50 -0.83 -2.48  0.02  0.57  0.73 -0.82 -0.54 -0.29  0.40  1.11  0.96
3 0.20  0.02  1.11 -0.27  0.09  0.81  1.21 -0.50  2.13  0.27  0.32  0.79  1.39
4 1.69 -0.35  0.49 -0.81  0.93  2.48  1.46 -0.98  1.03  1.18 -0.43  1.18  2.33
5 0.29  0.23  1.84  0.45  1.28  0.81  0.66  0.23  0.40 -0.32  0.36  0.45 -0.04
6 1.48 -0.52  0.30 -1.29  0.86  1.56  1.36 -0.18  0.66  0.73  0.40  0.34  2.23
```

How many principal components?

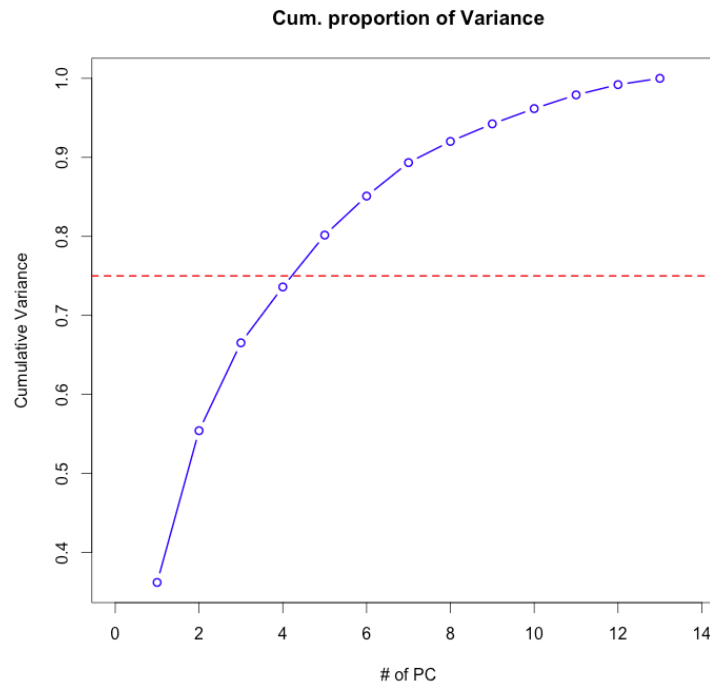
- Criteria #1: Slope change in variance



```
plot(c(1:13), (pca$sdev)^2, type="b", col="blue", lwd=2)  
abline(v=4, lwd=2, lty=2, col="red")
```

How many principal components?

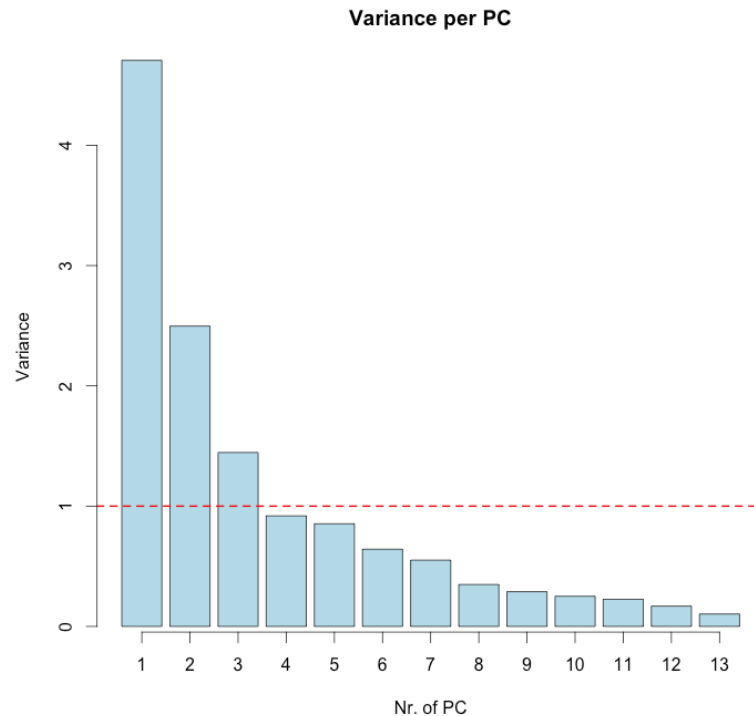
- Criteria #2: Cumulative variance



```
plot(cumsum((pca$sdev^2)/sum(pca$sdev^2)), type="b", lwd=2)  
abline(h=0.75, col="red", lty=2, lwd=2)
```

How many principal components?



- **Criteria #3: Variance > 1**



```
bp <- barplot(pca$sdev^2)
abline(h=1, col="red", lwd=2, lty=2)
```

PC1 vs. PC2 vs. PC4


PC1 vs. PC2



Note: Although PCA is very useful in finding the components with the highest variation, it does not always mean that the component are useful for separation in that order.

```
library(scatterplot3d)
pairs(pc$scores[,1:5], col=rainbow(3)[wine[,1]], asp=1)
scatterplot3d(pc$scores[,c(1,2,4)], color=rainbow(3)[wine[,1]])
```

Variables factor map (PCA)





Another example: cars

		Sports SUV Wagon Minivan Pickup AWD						RWD	Retail	Dealer	Engine	Cylinders	Horsepower	CityMPG	HighwayMPG	Weight	Wheelbase	Length	Width
Acura	3.5 RL	0	0	0	0	0	0	0	43755	39014	3.5	6	225	18	24	3880	115	197	72
Acura	MDX	0	1	0	0	0	1	0	36945	33337	3.5	6	265	17	23	4451	106	189	77
Acura	NSX S	1	0	0	0	0	0	1	89765	79978	3.2	6	290	17	24	3153	100	174	71
Acura	RSX	0	0	0	0	0	0	0	23820	21761	2.0	4	200	24	31	2778	101	172	68
Acura	TL	0	0	0	0	0	0	0	33195	30299	3.2	6	270	20	28	3575	108	186	72

```
> pca = prcomp(cars[,8:18], scale.=TRUE)
> round(pca$rotation[,1:2],2)
```

	PC1	PC2
Retail	-0.26	-0.47
Dealer	-0.26	-0.47
Engine	-0.35	0.02
Cylinders	-0.33	-0.08
Horsepower	-0.32	-0.29
CityMPG	0.31	0.00
HighwayMPG	0.31	0.01
Weight	-0.34	0.17
Wheelbase	-0.27	0.42
Length	-0.26	0.41
Width	-0.30	0.31



Cumulative Proportion

Recipes for interpretation

- **Coordinates**

- Use the bi-plot to summarize what each component means.

- **Correlations**

- For many datasets, the data clusters into groups of highly correlated attributes

- **Clusters**

- Clusters indicate a preference for particular combinations of attribute values. Summarize each cluster by its prototypical member.

- **Funnels**

- Funnels are wide at one end and narrow at the other. They happen when one dimension affects the variance of another, orthogonal dimension.

- **Voids**

- Voids are areas inside the range of the data which are unusually unpopulated.

